API Pricing

Contact sales

Pricing below reflects standard processing rates. To optimize cost and performance for different use cases, we also offer:

- **Batch API**: Save 50% on inputs and outputs with the Batch API and run tasks asynchronously over 24 hours.
- **Priority processing**: offers reliable, high-speed performance with the flexibility to pay-as-you-go.

Flagship models

Our frontier models designed to spend more time thinking before producing a response, making them ideal for complex, multi-step problems.

GPT-5

The best model for coding and agentic tasks across industries

Price

Cached input:
$0.125 / 1M tokens

Output:
$10.000 / 1M tokens

GPT-5 mini

A faster, cheaper version of GPT-5 for well-defined tasks

Price

Input:
$0.250 / 1M tokens

Cached input:
$0.025 / 1M tokens

Output:
$2.000 / 1M tokens

GPT-5 nano

The fastest, cheapest version of GPT-5—great for summarization and classification tasks

Price

Input:
$0.050 / 1M tokens

Cached input:
$0.005 / 1M tokens

Output:
$0.400 / 1M tokens

Fine-tuning our models

Customize our models to get even higher performance for your specific use cases.

Ask ChatGPT

GPT-4.1

Fine-tuning price

Input:
$3.00 / 1M tokens

Cached input:
$0.75 / 1M tokens

Output:
$12.00 / 1M tokens

Training:
$25.00 / 1M tokens

GPT-4.1 mini

Fine-tuning price

Input:
$0.80 / 1M tokens

Cached input:
$0.20 / 1M tokens

Output:
$3.20 / 1M tokens

Training:
$5.00 / 1M tokens

Fine-tuning price

Input:
$0.20 / 1M tokens

Cached input:
$0.05 / 1M tokens

Output:
$0.80 / 1M tokens

Training:
$1.50 / 1M tokens

o4-mini

Reinforcement fine-tuning price

Input:
$4.00 / 1M tokens

Cached input:
$1.00 / 1M tokens

Output:
$16.00 / 1M tokens

Training:
$100.00 / training hour

Explore detailed pricing ↗

Our APIs

Realtime API

Build low-latency, multimodal experiences including speech-to-speech.

Text −

GPT-4o

$5.00 / 1M input tokens	$2.50 / 1M cached input tokens	$20.00 / 1M output tokens

GPT-4o mini

$0.60 / 1M input tokens	$0.30 / 1M cached input tokens	$2.40 / 1M output tokens

Audio +

Image Generation API

Precise, high-fidelity image generation and editing with our latest multimodal model.

Text −

GPT-image-1

Image

GPT-image-1

$10.00 / 1M input tokens	$2.50 / 1M cached input tokens*	$40.00 / 1M output tokens

Prompts are billed similarly to other GPT models. Image outputs cost approximately $0.01 (low), $0.04 (medium), and $0.17 (high) for square images.
*available via the Responses API

For detailed token usage by image quality and size, see the docs.

Responses API

Our newest API combining the simplicity of Chat Completions with the built-in tool use of Assistants.

Price

Responses API is not priced separately. Tokens are billed at the chosen language model's input and output rates.

Chat Completions API

Build text-based conversational experiences.

Price

Assistants API

Build assistant-like experiences with our tools.

Price

Assistants API is not priced separately. Tokens are billed at the chosen language model's input and output rates.

Built-in tools

Extend model capabilities with built-in tools in the API Platform.

Code Interpreter
$0.03

File Search Storage
$0.10 / GB of vector storage per day (first GB free)

File Search Tool Call (Responses API only)
$2.50 / 1k tool calls

Web Search Tool Call
The tokens used for built-in tools are billed at the chosen model's per-token rates.
Web search content tokens: Search content tokens are tokens retrieved from the search index and fed to the model alongside your prompt to generate an answer. For gpt-4o and gpt-4.1 models, these tokens are included in the $25/1K calls cost. For o3 and o4-mini models, search content tokens are charged at model rate.

Models	Search Content	Cost
gpt-4o, gpt-4o-mini, gpt-4.1, and gpt-4.1-mini-models*	Search content tokens free	$25.00 / 1K calls

GPT-5, GPT-5-mini, GPT-5-nano, o3, o4-mini, o3-pro, and deep research models	Search content tokens billed at model rate	$10.00 / 1K calls

The billing dashboard will report gpt-4o and gpt-4o-mini search line items as 'web search tool calls | gpt-4o' and 'web search tool calls | gpt-4o-mini'

GB refers to binary gigabytes of storage (also known as gibibyte), where 1GB is 2^30 bytes.

Explore detailed pricing ↗

Explore our offerings for Enterprise customers: **Priority processing**, **Scale Tier** and **Reserved Capacity**.

FAQ

Which model should I use?

We recommend that developers use our large and mini GPT models for everyday tasks. Our large GPT models generally perform better on a wide range of tasks, while our mini GPT models are fast and inexpensive for simpler tasks.

Our large and mini reasoning models are ideal for complex, multi-step tasks and STEM use cases that require deep thinking about tough problems. You can choose the mini reasoning model if you're looking for a faster, more inexpensive option.

Do you offer an enterprise package or SLAs? +

Will I be charged for API usage in the Playground? +

How will I know how many tokens I've used each month? +

How can I manage my spending on the API platform? +

Is access to the API included in ChatGPT Plus, Team, Enterprise or Edu? +

How is pricing calculated for images? +

Start creating with OpenAI's powerful models.

Get started Contact sales >

GPT-4o

GPT-4o mini

Sora

Safety

Safety Approach

Security & Privacy

Trust & Transparency

Sora

Sora Overview

Features

Pricing

Sora log in ↗

API Platform

Platform Overview

Pricing

API log in ↗

Documentation ↗

Developer Forum ↗

Careers

Brand

Support

Help Center ↗

More

News

Stories

Livestreams

Podcast

OpenAI © 2015–2025

Manage Cookies

English United States